

ADITI MOKASHI

Pune, India | +91 77208 22204 | aditimokashi23@gmail.com | linkedin.com/in/aditi-mokashi23

PROFESSIONAL SUMMARY

Data Engineer & AI/ML Engineer with 2+ years of experience designing large-scale ETL/ELT pipelines, cloud data migrations, and production-grade GenAI/RAG systems across AWS and Azure. Experienced in Medallion Architecture, distributed data processing, and machine-learning-driven pipelines (embeddings, NLP, agentic AI), with a proven track record of taking ambiguous business problems to production systems operating at multi-terabyte, multi-million-record scale.

EXPERIENCE

Associate Consultant | Emergys Solutions, Pune, India | May 2025 – Present

- Developed the end-to-end migration of ~16TB of enterprise data (2TB structured, 14TB unstructured) from Remedyforce to AWS within a 1-month timeline, designing ETL pipelines using AWS Glue and AWS Glue Visual ETL for Bronze-to-Gold transformations.
 - Individually engineered the Bronze-to-Gold transformation pipelines, processing 10M+ records in under 30 minutes, with curated Gold-layer data persisted in PostgreSQL for downstream analytics and reporting.
- Architected an Agentic RAG chatbot for querying internal business logic and financial models, combining code embeddings, FAISS vector search, and LLMs for context-aware reasoning and answers to user queries; deployed as a production REST API.
- The chatbot was designed as a multi-agent routing architecture that dynamically directs queries to domain-specific agents, improving response relevance across logic-heavy enterprise use cases.
- Implemented LLM guardrails and multi-turn conversational memory, improving reliability of AI-generated response in production.
- Engineered a high-throughput data export pipeline (Python) using asynchronous streaming and disk-based writes, processing ~1M records in 15–20 minutes on constrained (7GB) infrastructure while supporting concurrent large-scale data extraction requests.
- Designed a pluggable data integration framework (Java, Spring Boot) connecting heterogeneous sources — Microsoft Fabric SQL Analytics, Azure Data Lake Storage, PostgreSQL — for automated metadata extraction, data governance, and business-rule-driven analysis.
- Built an orchestration pipeline implementing Medallion Architecture (Bronze/Silver/Gold) across Salesforce, RDBMS, and SharePoint sources on Azure Databricks (Genie, Unity Catalog, Ingestion Pipelines, Jobs), enabling real-time AI-powered querying on Gold-layer data.

Senior Software Engineer | Emergys Solutions, Pune, India | Jul 2024 – Apr 2025

- Built a configurable ETL tool using Python, FastAPI, and Azure App Service to ingest multiple CSV/Excel files and transform them into clean, standardized, application-ready data. Automated data cleaning, basic text/NLP normalization, data type and date standardization, geographic normalization to ISO codes (e.g., America/USA/US → USA), and currency conversion using transaction-date exchange rates.
- Built vendor intelligence orchestration pipelines (Azure Databricks, ADF) for batch processing of multiple files:
 - Normalization:** grouped duplicate vendor entries using Sentence Transformer embeddings, clustering, and prompt engineering, achieving ~95% accuracy across 10K vendors in ~17 minutes.
 - Classification:** categorized normalized vendors into predefined categories using Azure AI Agents with Bing grounding, achieving ~88% accuracy across 10K vendors.
- Developed automated reporting pipelines (Azure Functions, python-pptx) that generate power point presentations turning processed procurement and spend datasets into real-time business insights via an async, horizontally scalable service.

Software Engineering Intern | Emergys Solutions, Pune, India | Jan 2024 – Jun 2024

- Built a distributed data scraping and aggregation pipeline (Python) collecting and normalizing news data from multiple sources (Bing, Google, Yahoo) for downstream NLP processing.
- Developed a BERT-based NER model for domain-specific entity extraction (proof of concept).

PUBLICATIONS

Violation Eye – AI-Based Traffic Violation Reporting System | IEEE | July 2025

- Designed an AI-powered system combining computer vision, OCR, and real-time processing for traffic violation detection and reporting, as part of an internship with Equations Work.

SKILLS

Data Engineering & Big Data: AWS Glue, AWS Glue Visual ETL, Databricks, Azure Data Factory, ETL/ELT, Medallion Architecture (Bronze/Silver/Gold), Large-Scale Data Migration, Distributed Processing, Spark/PySpark, Data Governance, Unity Catalog

AI / ML & GenAI: RAG, Agentic AI, Multi-Agent Systems, LLM Guardrails, Prompt Engineering, Embeddings, Vector Search (FAISS), Sentence Transformers, BERT, NER, Azure OpenAI, Azure AI Agents

Languages & Frameworks: Python (FastAPI, Flask), Java (Spring Boot), SQL

Cloud Platforms: AWS (Glue), Azure (Databricks, Data Factory, Functions, App Service, Bing Grounding, AI Foundry, DevOps), Git

Databases: PostgreSQL, MySQL, SQL Server

EDUCATION

B.E. Computer Engineering | MMCOE, Pune | 2021 – 2024 | CGPA: 9.25

Diploma in Computer Engineering | Government Polytechnic, Pune | 2018 – 2021 | Percentage: 93.67